

# SUBMISSION OF WRITTEN WORK

Class code:

Name of course:

Course manager:

Course e-portfolio:

Thesis or project title:

Supervisor:

Full Name:

Birthdate (dd/mm-yyyy):

E-mail:

- |          |       |              |
|----------|-------|--------------|
| 1. _____ | _____ | _____@itu.dk |
| 2. _____ | _____ | _____@itu.dk |
| 3. _____ | _____ | _____@itu.dk |
| 4. _____ | _____ | _____@itu.dk |
| 5. _____ | _____ | _____@itu.dk |
| 6. _____ | _____ | _____@itu.dk |
| 7. _____ | _____ | _____@itu.dk |

# Contents

<b>1</b>	<b>Question 1</b>	<b>4</b>
1.1	.....	4
1.2	.....	4
1.3	.....	6
1.4	.....	7
1.5	.....	7
<b>2</b>	<b>Question 2</b>	<b>7</b>
2.1	.....	7
2.2	.....	8
2.3	.....	8
2.4	.....	8
2.5	.....	8
2.6	.....	9
<b>3</b>	<b>Question 3</b>	<b>10</b>
3.1	CPU pinning: .....	10
3.2	Load balancing: .....	10
<b>4</b>	<b>Question 4</b>	<b>10</b>
4.1	.....	10
4.2	.....	10
4.3	.....	10
<b>5</b>	<b>Question 5</b>	<b>11</b>
5.1	.....	12
5.2	.....	12
5.3	.....	13
5.3.1	Concurrency .....	13
5.3.2	Ordering .....	14
5.4	.....	15
5.5	.....	16
5.6	.....	16
5.7	.....	16
5.8	.....	17
<b>6</b>	<b>Question 6</b>	<b>17</b>
6.1	.....	17
6.2	.....	17
6.3	.....	17
<b>7</b>	<b>Question 7</b>	<b>17</b>
7.1	.....	17
7.2	.....	17
<b>8</b>	<b>Appendix</b>	<b>18</b>
8.1	TestQuickSelect.java .....	18
8.2	MyUnionFind.java .....	24

8.3	MyStack.java . . . . .	30
8.4	SecComSys.java . . . . .	32

# Practical Concurrent and Parallel Programming

Emil Lynegaard

December 12, 2017

*I hereby declare that I have answered the exam questions myself without any outside help.*

Throughout the report, I will be using inline code snippets, but include full files in the appendix. For all tests, the same machine will be used. Table 1 contains the results of **SystemInfo**:

---

OS	Linux; 4.13.12-1-ARCH; amd64
JVM	Oracle Corporation; 1.8.0_144
CPU	null; 8 "cores"
Date	2017-12-11T09:20:13+0100

Table 1: System Info

---

# 1 Question 1

## 1.1

Seeing as we are interested in seeing how well each implementation performs on random input of different sizes, we use Mark9 for the benchmarking, as it calculates the per element mean time and standard deviation.

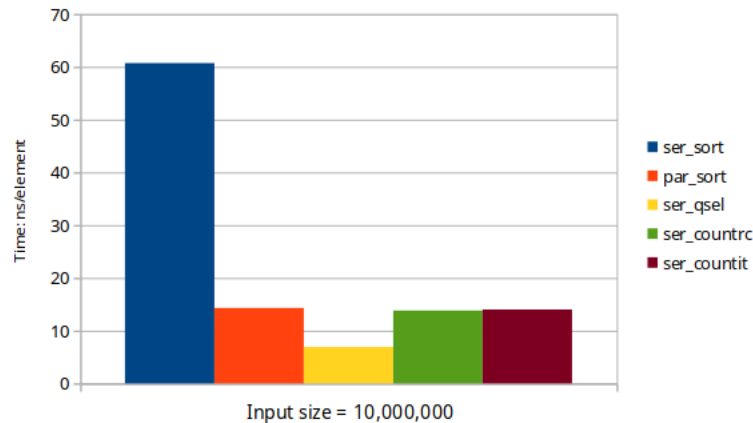


Figure 1: Plot of running times of given implementations. Tested with input size  $10^7$ .

From Graph 1 we see that, perhaps unsurprisingly, the serial `quickSelect` and `quickCount` implementations beat out serial sort entirely. We do however see parallel sort get rather close to the expected linear running time implementations, due to the test machine having a quad-core CPU, allowing an approximate 4 time speedup from the serial implementation. Between the expected linear running time algorithms, serial `quickSelect` beats out its `quickCount` counterparts. This may be due to `quickSelect` avoiding spending time allocating new memory, other than its initial copy.

## 1.2

Below is shown a parallel `quickCount` implementation using tasks, where `quickCountItTask` is the main function, and `filter` is the auxiliary method in which the parallel filtering is done. The filter method mainly serves to reduce code duplication as the original implementation is near identical for the filtering cases where we either have a too larger or a too small selected partition number.

To allow multiple tasks to safely increment `count`, we use a `final AtomicInteger` on line 28. In general for each iteration, all variables that we need to access from within the threads are made final, as we get compile errors otherwise. For task partitioning, we try to split up the input array in even intervals skipping the partition element. For undivisible numbers, we give the leftovers to the last created task. This can slow us down a little bit, since one task may end up doing more work than the others. In `filter` we take the same approach to partitioning, and now use an `AtomicInteger` on line 8, to safely keep track of what index in the new array `m` we

want to insert the next value in. This ensures that we never try to write to the same index of `m` twice.

Other than these parallelization measures, the overall structure of `quickCountItTask` is similar to that of the given `quickCountTask`.

---

```
1 // Takes an arr, a partition, the size of the output array and a BiFunction,
2 // returning an array of the given size containing elements from arr for which
3 // f.apply(arr[i], partition) returns true.
4 public static int[] filter(int[] arr, int partition, int size,
5     BiFunction<Integer,Integer,Boolean> f){
6     int[] m = new int[size];
7     ArrayList<Callable<Void>> filterers = new ArrayList<>();
8     final AtomicInteger j = new AtomicInteger(0);
9     final int step = arr.length/threadCount;
10    for(int i=0;i<threadCount;i++) {
11        final int from = i==0 ? 1 : i*step;
12        final int to = i==threadCount-1 ? arr.length : i*step+step;
13        filterers.add(() -> {
14            for(int h= from; h<to; h++)
15                if(f.apply(arr[h],partition)) m[j.getAndIncrement()]=arr[h];
16            return null;
17        });
18    }
19    try{ executor.invokeAll(filterers);
20    } catch (InterruptedException e) { System.err.println("Threads interrupted");}
21    return m;
22 }
23
24 static ExecutorService executor = Executors.newWorkStealingPool();
25 public static int quickCountItTask(int[] in) {
26     int target = in.length/2;
27     do {
28         final AtomicInteger count = new AtomicInteger(0);
29         final int[] inp = in;
30         final int n = inp.length, p = inp[0];
31         final int step = n/threadCount;
32
33         //Counting
34         ArrayList<Callable<Void>> counters = new ArrayList<>();
35         for(int i=0;i<threadCount;i++) {
36             final int from = i==0 ? 1 : i*step; //skip pivot
37             //for undivisible numbers, just let the last thread take a larger chunk
38             final int to = i==threadCount-1 ? inp.length : i*step+step;
39             counters.add(() -> {
40                 for(int j = from; j<to; j++)
41                     if(inp[j]<p) count.getAndIncrement();
42                 return null;
43             });
44         }
45         try{ executor.invokeAll(counters);
46         } catch (InterruptedException e) { System.err.println("Threads interrupted");}
47
48         if (count.get() == target) return p; //Terminated
```

```

48
49 //Filtering
50 boolean tooLargeP = count.get() > target;
51 int size = tooLargeP ? count.get() : n-count.get()-1;
52 if(tooLargeP) {
53     in = filter(inp, p, size, (x,y) -> x < y);
54 } else {
55     in = filter(inp, p, size, (x,y) -> x >= y);
56     target=target-count.get()-1;
57 }
58 } while( true );
59 }

```

---

### 1.3

Since the machine used for testing has 4 physical cores, this is the ideal amount of threads for the quickSelectIt. There is however only little difference between 4 and 8 since it supports hyperthreading<sup>1</sup>. Because of this, I have set the `threadCount` to 4, and tested with various input sizes. Figure 2 depicts the results.

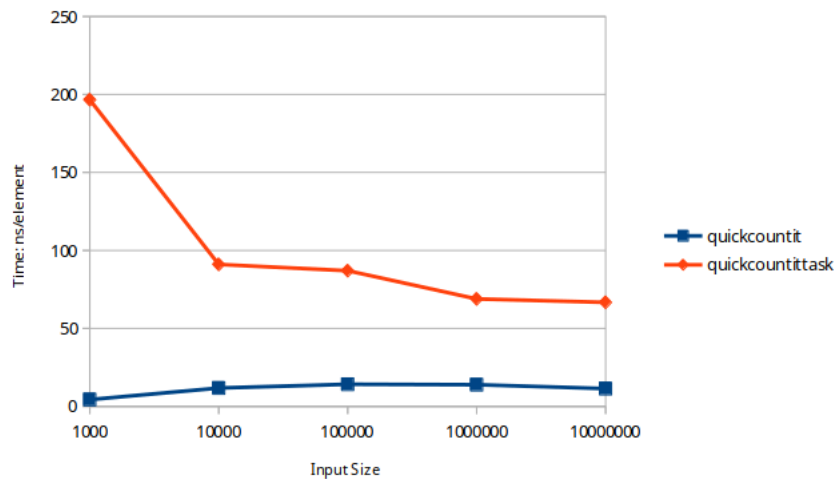


Figure 2: Plot of running times per element between `quickCountIt` and `quickCountItTask`

Perhaps in contrast to the expected outcome, `quickCountItTask` remains slower than `quickCountIt` even for large inputs. This can be due to how much it slows down on small inputs, with near  $200ns/element$  for inputs of size  $10^3$ . Based on the Figure 2 it seems that it will not be able to catch up to the serial implementation no matter the input size.

<sup>1</sup><https://ark.intel.com/products/80806/Intel-Core-i7-4790-Processor-8M-Cache-up-to-4.00-GHz>

## 1.4

We assume that a single partitioning step indicates one entire iteration of the `do-block` including counting and filtering. To test a single partitioning step we temporarily modify `quickCountItTask` to return after a single iteration. Table 2 depicts the running time per element for inputs of size  $10^4$  with a varying number of threads for a single step. Here we see that on smaller input sizes, the overhead of using multiple threads becomes increasingly noticeable.

Threads	1	2	4	8	16	32
Time: ns/element	15.0	28.4	32.7	35.7	35.9	36.4

Table 2: Running times per element for a single partitioning step in `quickCountItTask`, with input size  $10^4$ , and a varying thread count

Similarly we also try to pin the thread count to 4 and vary the input size. This is depicted in Table 3. Here we see that with the pinned thread count, we seem to scale positively in the size of the input up until  $10^6$ . This makes sense in terms of the benefits of multithreading outweighing the overhead, when each thread gets a sufficiently large part of the input to work with.

Input size	$10^3$	$10^4$	$10^5$	$10^6$	$10^7$	$10^8$
Time: ns/element	26.7	32.5	32.4	18.1	19.7	22.4

Table 3: Running times per element for a single partitioning step in `quickCountItTask`, with input size of thread count of 4 and varying input size

## 1.5

To make our implementation a hybrid, we declare a constant `CUTOFF` and add the following line to the top of the `do-block` in `quickCountItTask`.

---

```
1 if (in.length <= CUTOFF) return quickCountIt(in);
```

---

This change does make `quickCountItTask` faster, but even as a hybrid, it is unable to beat the sequential version, meaning that it still remains best to use the sequential version. This finding may either be due to the overhead of thread synchronization, or that the parallelization could be implemented in a more thread-friendly way.

## 2 Question 2

### 2.1

---

```
1 (int) Arrays.stream(inp).skip(1).filter(i -> i < p).count();
```

---

Where `inp` is our input array, and `p` is our current partition candidate. Skip the first element as this is the index of the partition element with which we do not wish to compare.



## 2.2

---

```
1     Arrays.stream(inp).skip(1).filter(i -> i < p).toArray();
2     Arrays.stream(inp).skip(1).filter(i -> i >= p).toArray();
```

---

## 2.3

---

```
1 public static int quickCountStream(int[] inp) {
2     int partition=-1, count=0, n=inp.length;
3     int target = n/2;
4     do {
5         partition=inp[0];
6         final int p = partition;
7         n=inp.length;
8         count = (int) Arrays.stream(inp).skip(1).filter(i -> i < p).count();
9         if (count == target) break;
10        if (count > target){
11            inp = Arrays.stream(inp).skip(1).filter(i -> i < p).toArray();
12        }else{
13            inp = Arrays.stream(inp).skip(1).filter(i -> i >= p).toArray();
14            target=target-count-1;
15        }
16    } while( true );
17    return partition; // we are on target
18 }
```

---

Combining the two we get above implementation, which yields correct results.

## 2.4

---

```
1     Arrays.stream(inp).parallel().skip(1).filter(i -> i < p).count();
2     Arrays.stream(inp).parallel().skip(1).filter(i -> i < p).toArray();
3     Arrays.stream(inp).parallel().skip(1).filter(i -> i >= p).toArray();
```

---

Here we simply throw `.parallel()` onto the pipelines from before.

## 2.5

---

```
1 public static int quickCountStream(int[] inp) {
2     int partition=-1;
3     int target = inp.length/2;
4     // Since we have to be working with boxed Integers. We start off by converting.
5     List<Integer> list = Arrays.stream(inp).boxed().collect(Collectors.toList());
6     do {
7         partition = list.get(0);
8         final Integer p = partition;
```

```

9      Map<Boolean, List<Integer>> res = list.stream().skip(1).parallel()
10         .collect(Collectors.partitioningBy(i -> i < p));
11
12      List<Integer> smaller = res.get(true);
13      List<Integer> bigger = res.get(false);
14
15      if (smaller.size() == target) break;
16      if (smaller.size() > target) list = smaller;
17      else {
18          target=target-smaller.size()-1;
19          list = bigger;
20      }
21  } while( true );
22  return partition; // we are on target
23 }

```

To avoid having to constantly do boxing, since `Collectors` does not work with primitives, we start off by converting our (`[] inp`) to `List<Integer>`.

The `partitioningBy` collector gives us a map of all with two entries. The `true` entry on line 12, holding all elements larger than our partition element, and false entry on line 13 holding all the elements larger than or equal to our partition element. The size of `smaller` now represents our count from before, and the remainder of the code is similar to the given code.

## 2.6

In Figure 3 we have tested all implementations from Figure 1 again, and added the hybrid implementation as well as the stream implementations. Here the `CUTOFF` for the parallel iterative `quickCount` is set to  $10^4$ . Based on these tests, it appears that every parallel implementation is slower than the serial `quickSelect` and the serial recursive and iterative `quickCount`. While this is perhaps somewhat surprising, we can deduce that we either need smarter ways of parallelizing the implementations, or that `quickCount` does not lend itself positively to parallelization on my machine, despite it being very reasonable to implement.

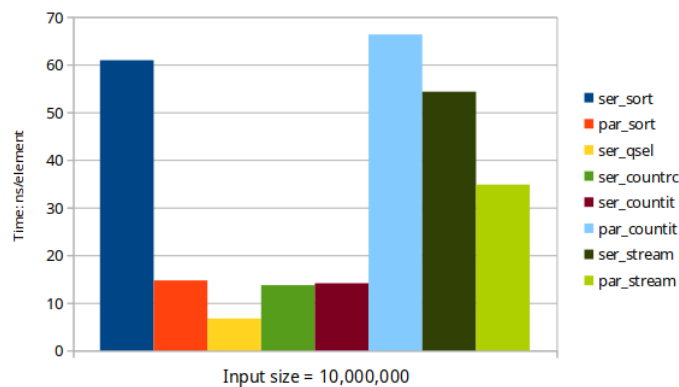


Figure 3: Plot of running times including parallel iterative and stream based. Tested with input size  $10^7$ .

## 3 Question 3

### 3.1 CPU pinning:

If we assume that the total number of threads is equal to the number of cores on the machine of execution, CPU pinning can be advantageous if the program gets all of the machine's CPU time, and the list is evenly partitioned between threads. In these specific conditions, not having to schedule threads may be beneficial as these conditions should cause various threads to finish their iteration's work simultaneously. On the other hand, if the input is not evenly distributed and the system is using resources on other tasks as well, one partition may end up being poorly scheduled on a core with less free CPU time than one of the other available cores. Another benefit of thread pinning, which may be more significant, is that this will allow the threads to keep most of the relevant values in their designated cores' L1 cache, leading to less cache misses<sup>2</sup>. This benefit may even extend to L2 cache for some CPUs<sup>3</sup>.

### 3.2 Load balancing:

By having each thread own certain data ranges, we may end up with several threads having nothing or very little to do after few iterations. For example, if the input is near sorted and we assign threads consecutive ranges, some threads may up filtering out all their elements in the first iteration depending on the selected pivot.

## 4 Question 4

### 4.1

Since `x` and `y` are parameters, we are not guaranteed to always grab locks in the same order, hence we are prone to dead-locks.

### 4.2

---

```
1 union(0,1)
2 union(1,0)
```

---

Above example when executed in parallel may lead to the first call grabbing the lock on `nodes[0]`, the second call grabbing the lock on `nodes[1]` and both calls thereafter waiting to get the lock on the node that their counterpart already grabbed.

### 4.3

We modify the given `concurrent()` method in class `UnionFindTest` in file `MyUnionFind.java` to target the issue we identified in 4.1.

---

<sup>2</sup>[https://en.wikipedia.org/wiki/CPU\\_cache#Cache\\_miss](https://en.wikipedia.org/wiki/CPU_cache#Cache_miss)

<sup>3</sup>[https://en.wikipedia.org/wiki/CPU\\_cache#Cache\\_hierarchy\\_in\\_a\\_modern\\_processor](https://en.wikipedia.org/wiki/CPU_cache#Cache_hierarchy_in_a_modern_processor)

---

```

1 public void deadlock(final int size, final UnionFind uf) throws Exception {
2     final int[] numbers = new int[size];
3     for (int i = 0; i < numbers.length; ++i) numbers[i] = i;
4     final int threadCount = 32;
5     final CyclicBarrier startBarrier = new CyclicBarrier(threadCount+1),
6         stopBarrier = startBarrier;
7     Collections.shuffle(Arrays.asList(numbers));
8     for (int i = 0; i < threadCount; ++i) {
9         final boolean reverse = i%2==0;
10        Thread ti = new Thread(new Runnable() { public void run() {
11            try { startBarrier.await(); } catch (Exception exn) { }
12            if (reverse)
13                for (int j=0; j<100; j++)
14                    for (int i = 0; i < numbers.length - 1; ++i)
15                        uf.union(numbers[i], numbers[i + 1]);
16            else
17                for (int j=0; j<100; j++)
18                    for (int i = 0; i < numbers.length - 1; ++i)
19                        uf.union(numbers[i + 1], numbers[i]);
20            try { stopBarrier.await(); } catch (Exception exn) { }
21        }});
22        ti.start();
23    }
24    startBarrier.await();
25    stopBarrier.await();
26    final int root = uf.find(0);
27    for (int i : numbers) {
28        assertEquals(uf.find(i), root);
29    }
30    System.out.println("No deadlocks");
31 }

```

---

As seen from line 10 to 20, half the threads will now be attempting to union nodes in reverse order of the other half. From my tests, calling the above defined `deadlock` method with parameters shown below, deadlocked every time.

---

```

1 UnionFindTest test = new UnionFindTest();
2 test.deadlock(itemCount, new BogusFineUnionFind(itemCount));

```

---

By merely adding a check to the given `union()` method in class `BogusFineUnionFind` to ensure we always lock the lowest entry of the `nodes` array first, the deadlock test executes without deadlocking.

## 5 Question 5

### Specifications:

1. `pop` returns an inserted item or the value null. It might block until another concurrent operation completes, but it will return without delay if no other operation is happening

simultaneously. In particular, it will not block until another thread inserts some element.

2. for each element that is pushed, there is at most one pop operation that returns that element.
3. If there are no further concurrent operations, pop will succeed (i.e. return a non-null value) if so far there have been more successful push than pop operations.
4. If processor A pushed two elements x and y in this order, and processor B pops both elements, then this happens in reverse order. (There is no further constraint on ordering).

## 5.1

---

```
1 import java.util.LinkedList;
2 public class MyStack<T> {
3     private Object lock;
4     private LinkedList<T> stack;
5
6     public MyStack(){
7         lock = new Object();
8         stack = new LinkedList<T>();
9     }
10
11     public void push(T obj) {
12         synchronized(lock){
13             stack.push(obj);
14         }
15     }
16
17     public T pop() {
18         synchronized(lock){
19             return stack.peek() != null ? stack.pop() : null;
20         }
21     }
22 }
```

---

Above generic implementation utilizes that Java's `LinkedList` ships with `push` and `pop`. Alternatively we could use the combination `addLast` and `removeLast` or `addFirst` and `removeFirst` of which the second pair by Java's documentation is equivalent to `push` and `pop`<sup>4</sup>. In `pop`, given that the specification states that we should return null if the list is empty, we use `peek` to check if there is a first element, if there isn't return null, otherwise `pop`. The locking could also be done implicitly on `this` by marking the methods as `synchronized`, but here we use an explicit global lock object as it seems more in line with specification number 1.

## 5.2

Given the initially stated specifications, below are added bullet points (letters) matching specifications, describing why the implementation from 5.1 is sufficient.

---

<sup>4</sup><https://docs.oracle.com/javase/7/docs/api/java/util/LinkedList.html>

1. (a) We take care of this explicitly in the implementation for `MyStack` in section 5.1, where we on line 19, `peek` prior to popping. If we blindly popped on an empty list, we would get a `NoSuchElementException`<sup>5</sup>.
2. (a) Given that `pop` removes the single first element, and `push` only adds the element once, `pop` may only return an element pushed once.  
 (b) Furthermore, since we are using a single lock, everything happens sequentially, removing any chance of reading off an element that was removed in a different thread.
3. (a) Same as 2.a.
4. (a) Since everything is handled sequentially due to the global lock, we have the guarantee that if one thread pushes two items, these will be pushed in the order of which the thread called push.  
 (b) Given 3.a. and Java's Documentation stating that `push` and `pop` adds/removes from the head of the list, elements are bound to be popped in reverse compared to the order of which they were pushed.

## 5.3

### 5.3.1 Concurrency

To test for everything except reverse ordering, we define the below shown test `concurrentTest`.

---

```

1 public static void concurrentTest(final int size, final int threads,
2     MyStack<Integer> stack) throws Exception {
3     final CyclicBarrier startBarrier = new CyclicBarrier(threads+1),
4         stopBarrier = startBarrier;
5
6     final int range = size/threads;
7     for (int i = 0; i < threads; ++i) {
8         final int nr = i;
9         Thread ti = new Thread(new Runnable() { public void run() {
10             try { startBarrier.await(); } catch (Exception exn) { }
11             for(int j = range*nr; j<range*nr+range; j++)
12                 stack.push(j);
13             try { stopBarrier.await(); } catch (Exception exn) { }
14             });
15         ti.start();
16     }
17     startBarrier.await();
18     stopBarrier.await();
19     startBarrier.reset();
20     stopBarrier.reset();
21
22     final Set<Integer> pops = ConcurrentHashMap.newKeySet();
23     for (int i = 0; i < threads; ++i) {
24         final int nr = i;
25         Thread ti = new Thread(new Runnable() { public void run() {

```

---

<sup>5</sup>[https://docs.oracle.com/javase/7/docs/api/java/util/LinkedList.html#pop\(\)](https://docs.oracle.com/javase/7/docs/api/java/util/LinkedList.html#pop())

```

26         try { startBarrier.await(); } catch (Exception exn) { }
27         for(int j = range*nr; j<range*nr+range; j++)
28             pops.add(stack.pop());
29         try { stopBarrier.await(); } catch (Exception exn) { }
30     }));
31     ti.start();
32 }
33
34 startBarrier.await();
35 stopBarrier.await();
36
37 if (pops.size() == size) System.out.println("Concurrency works :");
38 else System.out.println("Concurrency doesn't work :(");
39 }

```

---

We create `CyclicBarriers` on line 3, which we use to ensure that our threads are working simultaneously. We then assign different threads ranges between 0 and `size` in which they will push all numbers to the stack.

When all the threads are ready, we call `startBarrier.await()` on line 17 to start the pushing and `stopBarrier.await()` on the next line to wait for all of them to finish pushing.

We now divide the work similarly for popping threads. For each number they pop, we add it to the `ConcurrentHashMap` backed `Set` defined on line 22. When all threads are done popping, we can check whether the size of our `Set` is equal to the number of items we tried to add. Since there are no duplicates in a `Set`, and we only pushed unique numbers, if these are equal we can with reasonably confidence say that our implementation is working.

**Scaling:** In terms of scaling, since `MyStack` functions entirely sequentially, it scales poorly with multiple threads. In fact, the more threads we use, the more time will be spent locking and waiting for locks, making it slower and slower. In table 4 behavior is illustrated.

Threads	1	2	4	8	16	32
Time (sec)	7.363	8.159	8.203	8.413	10.071	10.832

Table 4: Test times for `concurrentTest` with size  $10^7$

### 5.3.2 Ordering

To test that order works for multiple threads, we push number  $[0..n]$  onto the stack from thread A, and pop  $n$  items off the stack from thread B, checking that these are the numbers  $[n..0]$ . This is shown below in method `testOrder`.

---

```

1 public static void testOrder(int n, MyStack<Integer> stack){
2     final AtomicBoolean working = new AtomicBoolean(true);
3     Thread A = new Thread(new Runnable() { public void run() {
4         for(int i=0; i<n; i++) stack.push(i);
5     }});
6     Thread B = new Thread(new Runnable() { public void run() {
7         for(int i=0; i<n; i++)
8             working.compareAndSet(true, n-1-i == stack.pop());

```

```

9      });
10     try {
11         A.start(); A.join();
12         B.join(); B.join();
13     } catch (Exception e) {
14         System.out.println("Order dies >:(");
15     }
16     if(working.get()) System.out.println("Order works :)");
17     else System.out.println("Order doesn't work :(");
18 }

```

---

This behaviour was untested in `concurrentTest`, so a separate straight forward test to clear this up was needed.

## 5.4

Below is shown an implementation using striping, with a total of 32 stripes. To determine the stripe use `Thread.currentThread().hashCode()%STRIPEs` as shown on line 16 and 24. We use an `ArrayList` to store our `LinkedLists` since we cannot create arrays of parameterized types<sup>6</sup>.

In `pop`, we use `i%STRIPEs` to iterate through the stacks, starting from the stack at the computed stripe, with wraparound.

---

```

1  import java.lang.*;
2  import java.util.*;
3  public class MyStack<T> {
4      private Object lock;
5      private final List<LinkedList<T>> stacks;
6      private static final int STRIPES = 32;
7
8      public MyStack(){
9          lock = new Object();
10         stacks = new ArrayList<LinkedList<T>>();
11         for(int i = 0; i < STRIPES; i++)
12             stacks.add(new LinkedList<T>());
13     }
14
15     public void push(T obj) {
16         int stripe = Thread.currentThread().hashCode()%STRIPEs;
17         LinkedList<T> stack = stacks.get(stripe);
18         synchronized(stack){
19             stack.push(obj);
20         }
21     }
22
23     public T pop() {
24         int stripe = Thread.currentThread().hashCode()%STRIPEs;
25         for (int i = stripe; i < stripe+STRIPEs; i++){
26             LinkedList<T> stack = stacks.get(i%STRIPEs);

```

---

<sup>6</sup><https://docs.oracle.com/javase/tutorial/java/generics/restrictions.html#createArrays>



```

27         synchronized(stack){
28             if(stack.size() == 0) continue;
29             return stack.pop();
30         }
31     }
32     return null;
33 }
34 }

```

---

## 5.5

With the new striping such that we are actually running concurrently we see improved performance in our tests. The fastest execution, as seen in table 5 was with 16 threads where it ran in 5.866 seconds. This is in contrast to the single threaded execution from table 4 that ran in 7.363 seconds with a single thread. Considering that we have 4 physical cores available, this is not that big of a performance improvement. This could be due to many threads hashing to the same stripe or due to the added iteration through all the stacks when we run into an empty one.

Threads	1	2	4	8	16	32
Time (sec)	8.016	7.473	7.465	6.313	5.866	7.337

Table 5: Test times for MyStack with striping for concurrentTest with size  $10^7$

## 5.6

Given our implementation described in 5.4, we modify line 29 to say:

---

```

1 return i == stripe ? stack.pop() : stack.removeLast();

```

---

If we are on our own stripe's stack, pop from the front, otherwise pop from the back.

## 5.7

Below is a description in pseudocode of what will go wrong given our changes in 5.6.

---

```

1 //Thread A // Stripe 0
2 myStack.push(0);
3 myStack.push(1);
4
5 //Thread B // Stripe 1
6 myStack.pop(); // this will return 0 - should be 1
7 myStack.pop(); // this will return 1 - should be 0

```

---

## 5.8

The code described in section 5.3.2 detects this issue and outputs "Order doesn't work :(" as expected.

## 6 Question 6

### 6.1

The given code has the flaw that it allows multiple threads to get into the else block before anyone changes the `state`. This means that in an example with thread A calling `consensus(x)` and thread B calling `consensus(y)`, both may retrieve return values indicating that parameter value `x` or `y` is the consensus, whereas only one of them will be stored in `state`. Hence we have a race condition.

### 6.2

Using synchronization on some lock, in this case on `this`, has a problem with termination. If one process is to fail while holding the lock, or fall asleep for an extended amount of time, the whole system will halt as they wait to retrieve the lock. As such, this implementation is not fault tolerant.

### 6.3

Firstly, assuming that we're talking Java, this would fail to typecheck, as we're returning an `AtomicInteger` from a method that supposedly returns an `int` and there is no implicit conversion from `AtomicInteger` to `int`. Even with this fixed, it would still fail, since two threads may enter the while loop before state changes, causing one of them to be stuck in the loop forever, since `state.compareAndSet(-1,x)` will then be returning false over and over. This will work if only one thread ever gets to enter the while loop and successfully `compareAndSets` the `state`, since all other threads then merely get the value of the state, leading to consensus.

## 7 Question 7

### 7.1

Implementation found in appendix section 8.4.

### 7.2

Below is shown the output of three distinct runs of the Secure Communication System with Java+Akka.

```
$ java -cp scala.jar:akka-actor.jar:akka-config.jar:. SecComSys
public key: 18
private key: 8
cleartext: 'SECRET'
encrypted: 'KWUJWL'
decrypted: 'SECRET'
```

```
$ java -cp scala.jar:akka-actor.jar:akka-config.jar:. SecComSys
public key: 2
private key: 24
cleartext: 'SECRET'
encrypted: 'UGETGV'
decrypted: 'SECRET'
```

```
$ java -cp scala.jar:akka-actor.jar:akka-config.jar:. SecComSys
public key: 5
private key: 21
cleartext: 'SECRET'
encrypted: 'XJHWJY'
decrypted: 'SECRET'
```

## 8 Appendix

### 8.1 TestQuickSelect.java

---

```
1  import java.util.Arrays;
2  import java.util.Collections;
3  import java.util.Random;
4  import java.util.ArrayList;
5  import java.util.HashSet;
6  import java.util.List;
7  import java.util.Map;
8  import java.util.Comparator;
9  import java.util.stream.Collectors;
10 import java.util.stream.IntStream;
11 import java.util.stream.Stream;
12 import java.util.function.IntFunction;
13 import java.util.function.IntToDoubleFunction;
14 import java.util.function.Function;
15 import java.util.concurrent.*;
16 import java.util.concurrent.atomic.*;
17 import java.util.function.BiFunction;
18
19 class TestQuickSelect {
20     public static int medianSort(int[] inp) {
21         int w[] = Arrays.copyOf(inp, inp.length);
22         Arrays.sort(w);
23         return w[w.length/2];
24     }
}
```

```

25 public static int medianPSort(int[] inp) {
26     int w[] = Arrays.copyOf(inp, inp.length);
27     Arrays.parallelSort(w);
28     return w[w.length/2];
29 }
30
31 public static int partition(int[] w, int min, int max) {
32     int p = min; // use w[p] as pivot
33     int left=min+1, right = max-1;
34     while(left <= right) {
35         while( w[left] <= w[p] && left < right ) left++;
36         while( w[right] > w[p] && left <= right ) right--;
37         if(left >= right) break;
38         int t=w[left]; w[left]=w[right]; w[right]=t;
39     }
40     int t=w[p]; w[p]=w[right]; w[right]=t;
41     return right;
42 }
43
44 public static int quickSelect(int[] inp) {
45     int w[] = Arrays.copyOf(inp, inp.length);
46     return quickSelect(w,0,w.length,w.length/2);
47 }
48 public static int quickSelect(int[] w, int min, int max, int target) {
49     int p = partition(w,min,max);
50     if( p < target ) return quickSelect(w,p+1,max,target);
51     if( p > target ) return quickSelect(w,min,p,target);
52     return w[target]; // p==target
53 }
54
55 public static int quickSelectIt(int[] inp) {
56     int w[] = Arrays.copyOf(inp, inp.length);
57     int target = w.length/2;
58     int p = -1, min=0, max=w.length;
59     do{
60         p = partition(w,min,max);
61         if( p < target ) min=p+1;
62         if( p > target ) max=p;
63         // System.out.println(" "+p+" "+target);
64     } while(p!=target);
65     return w[p];
66 }
67
68 public static int quickCountRec(int[] inp, int target) {
69     final int p=inp[0], n=inp.length;
70     int count=0;
71     for(int i=1;i<n;i++) if(inp[i]<p) count++;
72     if(count > target) {
73         int m[] = new int[count];
74         int j=0;
75         for(int i=1;i<n;i++) if(inp[i]<p) m[j++]=inp[i];
76         return quickCountRec(m,target);
77     }

```

```

78     if(count < target) {
79         int m[] = new int[n-count-1];
80         int j=0;
81         for(int i=1;i<n;i++) if(inp[i]>=p) m[j++]=inp[i];
82         return quickCountRec(m,target-count-1);
83     }
84     return p; // we are on target
85 }
86
87 public static int quickCountIt(int[] inp) {
88     int p=-1, count=0, n=inp.length;
89     int target = n/2;
90     do {
91         p=inp[0];
92         count=0;
93         n=inp.length;
94         for(int i=1;i<n;i++) if(inp[i]<p) count++;
95         if(count > target) {
96             int m[] = new int[count];
97             int j=0;
98             for(int i=1;i<n;i++) if(inp[i]<p) m[j++]=inp[i];
99             inp = m;
100             continue;
101         }
102         if(count < target) {
103             int m[] = new int[n-count-1];
104             int j=0;
105             for(int i=1;i<n;i++) if(inp[i]>=p) m[j++]=inp[i];
106             inp =m;
107             target=target-count-1;
108             continue;
109         }
110         break;
111     } while( true );
112     return p; // we are on target
113 }
114
115 // Takes an arr, a partition, the size of the output array and a BiFunction,
116 // returning an array of the given size containing elements of arr for which
117 // f.apply(arr[i], partition) returns true.
118 public static int[] filter(int[] arr, int partition, int size,
119     BiFunction<Integer,Integer,Boolean> f){
120     int[] m = new int[size];
121     ArrayList<Callable<Void>> filterers = new ArrayList<>();
122     final AtomicInteger j = new AtomicInteger(0);
123     final int step = arr.length/threadCount;
124     for(int i=0;i<threadCount;i++) {
125         final int from = i==0 ? 1 : i*step;
126         final int to = i==threadCount-1 ? arr.length : i*step+step;
127         filterers.add(() -> {
128             for(int h= from; h<to; h++)
129                 if(f.apply(arr[h],partition)) m[j.getAndIncrement()]=arr[h];
130             return null;
131         });
132     }
133     for(Callable<Void> filterer : filterers) filterer.call();
134     return m;
135 }

```

```

130     });
131 }
132 try{ executor.invokeAll(filterers);
133 } catch (InterruptedException e) { System.err.println("Threads interrupted");}
134 return m;
135 }
136
137 final static ExecutorService executor = Executors.newWorkStealingPool();
138 final static int CUTOFF = 10_000;
139 public static int quickCountItTask(int[] in) {
140     int target = in.length/2;
141     do {
142         if (in.length <= CUTOFF) return quickCountIt(in);
143         final AtomicInteger count = new AtomicInteger(0);
144         final int[] inp = in;
145         final int n = inp.length, p = inp[0];
146         final int step = n/threadCount;
147
148         //Counting
149         ArrayList<Callable<Void>> counters = new ArrayList<>();
150         for(int i=0;i<threadCount;i++) {
151             final int from = i==0 ? 1 : i*step; //skip pivot
152             // for undivisible numbers, just let the last thread take a larger chunk
153             final int to = i==threadCount-1 ? inp.length : i*step+step;
154             counters.add(() -> {
155                 for(int j= from; j<to; j++)
156                     if(inp[j]<p) count.getAndIncrement();
157                 return null;
158             });
159         }
160         try{ executor.invokeAll(counters);
161         } catch (InterruptedException e) { System.err.println("Threads
162             interrupted");}
163
164         if (count.get() == target) return p; //Terminated
165
166         //Filtering
167         boolean tooLargeP = count.get() > target;
168         int size = tooLargeP ? count.get() : n-count.get()-1;
169         if(tooLargeP) {
170             in = filter(inp, p, size, (x,y) -> x < y);
171         } else {
172             in = filter(inp, p, size, (x,y) -> x >= y);
173             target=target-count.get()-1;
174         }
175     } while( true );
176 }
177
178 public static int quickCountStreamP(int[] inp) {
179     int partition=-1;
180     int target = inp.length/2;
181     // Since we have to be working with boxed Integers. We start off by converting.
182     List<Integer> list = Arrays.stream(inp).boxed().collect(Collectors.toList());

```

```

182     do {
183         partition = list.get(0);
184         final Integer p = partition;
185         Map<Boolean, List<Integer>> res = list.stream().skip(1).parallel()
186             .collect(Collectors.partitioningBy(i -> i < p));
187
188         List<Integer> smaller = res.get(true);
189         List<Integer> bigger = res.get(false);
190
191         if (smaller.size() == target) break;
192         if (smaller.size() > target) list = smaller;
193         else {
194             target=target-smaller.size()-1;
195             list = bigger;
196         }
197     } while( true );
198     return partition; // we are on target
199 }
200
201 public static int quickCountStream(int[] inp) {
202     int partition=-1;
203     int target = inp.length/2;
204     // Since we have to be working with boxed Integers. We start off by converting.
205     List<Integer> list = Arrays.stream(inp).boxed().collect(Collectors.toList());
206     do {
207         partition = list.get(0);
208         final Integer p = partition;
209         Map<Boolean, List<Integer>> res = list.stream().skip(1)
210             .collect(Collectors.partitioningBy(i -> i < p));
211
212         List<Integer> smaller = res.get(true);
213         List<Integer> bigger = res.get(false);
214
215         if (smaller.size() == target) break;
216         if (smaller.size() > target) list = smaller;
217         else {
218             target=target-smaller.size()-1;
219             list = bigger;
220         }
221     } while( true );
222     return partition; // we are on target
223 }
224
225 public static final int threadCount = 4;
226 public static void main( String [] args ) {
227     SystemInfo();
228     int a[] = new int[Integer.parseInt(args[0])];
229     Random rnd = new Random();
230     if( args.length == 1 ) {
231         int nrIt = 10;
232         for(int ll=0;ll<nrIt;ll++) {
233             rnd.setSeed(23434+ll); // seed
234             for(int i=0;i<a.length;i++) a[i] = rnd.nextInt(4*a.length);

```

```

235         final int ra = quickCountRec(a,a.length/2);
236         final int rb = medianPSort(a);
237         if( ra !=rb ) {
238             System.out.println(ll);
239             System.out.println(ra);
240             System.out.println(rb);
241             System.exit(0);
242         }
243     }
244     System.out.println();
245 } else {
246     rnd.setSeed(23434+Integer.parseInt(args[1])); // seed
247     for(int i=0;i<a.length;i++) a[i] = rnd.nextInt(4*a.length);
248     System.out.println(medianPSort(a));
249     System.out.println(quickCountRec(a,a.length/2));
250 }
251 // System.exit(0);
252 int[] testArray = new int[]{9,2,4,3,5,7,1,8,9,6};
253 double d=0.0;
254 d += Mark9("serial sort", a.length, x -> medianSort(a));
255 d += Mark9("parall sort", a.length, x -> medianPSort(a));
256 d += Mark9("serial qsel", a.length, x -> quickSelect(a));
257 d += Mark9("ser countRc", a.length,x -> quickCountRec(a,a.length/2));
258 d += Mark9("ser countIt", a.length,x -> quickCountIt(a));
259 d += Mark9("par countIt", a.length,x -> quickCountItTask(a));
260 d += Mark9("countStream", a.length,x -> quickCountStream(a));
261 d += Mark9("countStreamP", a.length,x -> quickCountStreamP(a));
262 // d += Mark9("task countR", a.length,x -> quickCountRecTask(a,a.length/2));
263 System.out.println(d);
264 }
265
266 public static double Mark7(String msg, IntToDoubleFunction f) {
267     int n = 10, count = 1, totalCount = 0;
268     double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
269     do {
270         count *= 2;
271         st = sst = 0.0;
272         for (int j=0; j<n; j++) {
273             Timer t = new Timer();
274             for (int i=0; i<count; i++)
275                 dummy += f.applyAsDouble(i);
276             runningTime = t.check();
277             double time = runningTime * 1e9 / count;
278             st += time;
279             sst += time * time;
280             totalCount += count;
281         }
282     } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
283     double mean = st/n, sdev = Math.sqrt((sst - mean*mean*n)/(n-1));
284     System.out.printf("%-25s %15.1f ns %10.2f %10d\n", msg, mean, sdev, count);
285     return dummy / totalCount;
286 }
287

```



```

288 public static double Mark9(String msg, int size, IntToDoubleFunction f) {
289     int n = 5, count = 1, totalCount = 0;
290     double dummy = 0.0, runningTime = 0.0, st = 0.0, sst = 0.0;
291     do {
292         count *= 2;
293         st = sst = 0.0;
294         for (int j=0; j<n; j++) {
295             Timer t = new Timer();
296             for (int i=0; i<count; i++)
297                 dummy += f.applyAsDouble(i);
298             runningTime = t.check();
299             double time = runningTime * 1e9 / count; // microseconds
300             st += time;
301             sst += time * time;
302             totalCount += count;
303         }
304     } while (runningTime < 0.25 && count < Integer.MAX_VALUE/2);
305     double mean = st/n/size, sdev = Math.sqrt((sst - mean*mean*n)/(n-1))/size;
306     System.out.printf("%-25s %15.1f ns %10.2f %10d%n", msg, mean, sdev, count);
307     return dummy / totalCount;
308 }
309
310 public static void SystemInfo() {
311     System.out.printf("# OS: %s; %s; %s%n",
312         System.getProperty("os.name"),
313         System.getProperty("os.version"),
314         System.getProperty("os.arch"));
315     System.out.printf("# JVM: %s; %s%n",
316         System.getProperty("java.vendor"),
317         System.getProperty("java.version"));
318     // The processor identifier works only on MS Windows:
319     System.out.printf("# CPU: %s; %d \"cores\"%n",
320         System.getenv("PROCESSOR_IDENTIFIER"),
321         Runtime.getRuntime().availableProcessors());
322     java.util.Date now = new java.util.Date();
323     System.out.printf("# Date: %s%n",
324         new java.text.SimpleDateFormat("yyyy-MM-dd'T'HH:mm:ssZ").format(now));
325 }
326 }

```

---

## 8.2 MyUnionFind.java

```

1 import java.util.concurrent.atomic.AtomicInteger;
2 import java.util.concurrent.atomic.AtomicReferenceArray;
3
4 import java.util.concurrent.Callable;
5 import java.util.concurrent.CyclicBarrier;
6
7 import java.util.Arrays;
8 import java.util.Collections;
9

```

```

10 public class MyUnionFind {
11     public static void main(String[] args) throws Exception {
12         final int itemCount = 10_000;
13         {
14             UnionFindTest test = new UnionFindTest();
15             test.sequential(new FineUnionFind(5));
16             test.concurrent(itemCount, new FineUnionFind(itemCount));
17             test.deadlock(itemCount, new FineUnionFind(itemCount));
18         }
19         // Question 4.3
20         {
21             UnionFindTest test = new UnionFindTest();
22             test.sequential(new BogusFineUnionFind(5));
23             test.concurrent(itemCount, new BogusFineUnionFind(itemCount));
24             test.deadlock(itemCount, new BogusFineUnionFind(itemCount));
25         }
26     }
27 }
28
29 interface UnionFind {
30     int find(int x);
31     void union(int x, int y);
32     boolean sameSet(int x, int y);
33 }
34
35 // Test of union-find data structures, adapted from Florian Biermann's
36 // MSc thesis, ITU 2014
37
38 class UnionFindTest extends Tests {
39
40     public void sequential(UnionFind uf) throws Exception {
41         System.out.printf("Testing %s ... ", uf.getClass());
42         // Find
43         assertEquals(uf.find(0), 0);
44         assertEquals(uf.find(1), 1);
45         assertEquals(uf.find(2), 2);
46         // Union
47         uf.union(1, 2);
48         assertEquals(uf.find(1), uf.find(2));
49
50         uf.union(2, 3);
51         assertEquals(uf.find(1), uf.find(2));
52         assertEquals(uf.find(1), uf.find(3));
53         assertEquals(uf.find(2), uf.find(3));
54
55         uf.union(1, 4);
56         assertEquals(uf.find(1), uf.find(2));
57         assertEquals(uf.find(1), uf.find(3));
58         assertEquals(uf.find(2), uf.find(3));
59         assertEquals(uf.find(1), uf.find(4));
60         assertEquals(uf.find(2), uf.find(4));
61         assertEquals(uf.find(3), uf.find(4));
62     }

```

```

63
64 public void deadlock(final int size, final UnionFind uf) throws Exception {
65     final int[] numbers = new int[size];
66     for (int i = 0; i < numbers.length; ++i) numbers[i] = i;
67     // Populate threads
68     final int threadCount = 32;
69     final CyclicBarrier startBarrier = new CyclicBarrier(threadCount+1),
70         stopBarrier = startBarrier;
71     Collections.shuffle(Arrays.asList(numbers));
72     for (int i = 0; i < threadCount; ++i) {
73         final boolean reverse = i%2==0;
74         Thread ti = new Thread(new Runnable() { public void run() {
75             try { startBarrier.await(); } catch (Exception exn) { }
76             if (reverse)
77                 for (int j=0; j<100; j++)
78                     for (int i = 0; i < numbers.length - 1; ++i)
79                         uf.union(numbers[i], numbers[i + 1]);
80             else
81                 for (int j=0; j<100; j++)
82                     for (int i = 0; i < numbers.length - 1; ++i)
83                         uf.union(numbers[i + 1], numbers[i]);
84             try { stopBarrier.await(); } catch (Exception exn) { }
85             }}};
86         ti.start();
87     }
88     startBarrier.await();
89     stopBarrier.await();
90     final int root = uf.find(0);
91     for (int i : numbers) {
92         assertEquals(uf.find(i), root);
93     }
94     System.out.println("No deadlocks");
95 }
96
97 public void concurrent(final int size, final UnionFind uf) throws Exception {
98     final int[] numbers = new int[size];
99     for (int i = 0; i < numbers.length; ++i)
100         numbers[i] = i;
101     // Populate threads
102     final int threadCount = 32;
103     final CyclicBarrier startBarrier = new CyclicBarrier(threadCount+1),
104         stopBarrier = startBarrier;
105     Collections.shuffle(Arrays.asList(numbers));
106     for (int i = 0; i < threadCount; ++i) {
107         Thread ti = new Thread(new Runnable() { public void run() {
108             try { startBarrier.await(); } catch (Exception exn) { }
109             for (int j=0; j<100; j++)
110                 for (int i = 0; i < numbers.length - 1; ++i)
111                     uf.union(numbers[i], numbers[i + 1]);
112             try { stopBarrier.await(); } catch (Exception exn) { }
113             }}};
114         ti.start();
115     }

```

```

116         startBarrier.await();
117         stopBarrier.await();
118         final int root = uf.find(0);
119         for (int i : numbers) {
120             assertEquals(uf.find(i), root);
121         }
122         System.out.println("passed");
123     }
124 }
125
126 class Tests {
127     public static void assertEquals(int x, int y) throws Exception {
128         if (x != y)
129             throw new Exception(String.format("ERROR: %d not equal to %d%n", x, y));
130     }
131
132     public static void assertTrue(boolean b) throws Exception {
133         if (!b)
134             throw new Exception(String.format("ERROR: assertTrue"));
135     }
136 }
137 // Fine-locking union-find. Union and sameSet lock on the intrinsic
138 // locks of the two root Nodes involved. Find is wait-free, takes no
139 // locks, and performs no compression.
140
141 // The nodes[] array entries are never updated after initialization
142 // inside the constructor, so no need to worry about their visibility.
143 // But the fields of Node objects are written (by union and compress
144 // while holding locks), and read by find without holding locks, so
145 // must be made volatile.
146
147 class FineUnionFind implements UnionFind {
148     private final Node[] nodes;
149
150     public FineUnionFind(int count) {
151         this.nodes = new Node[count];
152         for (int x=0; x<count; x++)
153             nodes[x] = new Node(x);
154     }
155
156     public int find(int x) {
157         while (nodes[x].next != x)
158             x = nodes[x].next;
159         return x;
160     }
161
162     public void union(final int x, final int y) {
163         while (true) {
164             int rx = find(x), ry = find(y);
165             if (rx == ry)
166                 return;
167             else if (rx > ry) {
168                 int tmp = rx; rx = ry; ry = tmp;

```

```

169     }
170     // Now rx < ry; take locks in consistent order
171     synchronized (nodes[rx]) {
172         synchronized (nodes[ry]) {
173             // Check rx, ry are still roots, else restart
174             if (nodes[rx].next != rx || nodes[ry].next != ry)
175                 continue;
176             if (nodes[rx].rank > nodes[ry].rank) {
177                 int tmp = rx; rx = ry; ry = tmp;
178             }
179             // Now nodes[rx].rank <= nodes[ry].rank
180             nodes[rx].next = ry;
181             if (nodes[rx].rank == nodes[ry].rank)
182                 nodes[ry].rank++;
183             compress(x, ry);
184             compress(y, ry);
185         } }
186     }
187 }
188
189 // Assumes lock is held on nodes[root]
190 private void compress(int x, final int root) {
191     while (nodes[x].next != x) {
192         int next = nodes[x].next;
193         nodes[x].next = root;
194         x = next;
195     }
196 }
197
198 public boolean sameSet(int x, int y) {
199     return find(x) == find(y);
200 }
201
202 class Node {
203     private volatile int next, rank;
204
205     public Node(int next) {
206         this.next = next;
207     }
208 }
209 }
210
211 // Bogus Fine-locking union-find. Union and sameSet lock on the intrinsic
212 // locks of the two root Nodes involved. Find is wait-free, takes no
213 // locks, and performs no compression.
214
215 // The nodes[] array entries are never updated after initialization
216 // inside the constructor, so no need to worry about their visibility.
217 // But the fields of Node objects are written (by union and compress
218 // while holding locks), and read by find without holding locks, so
219 // must be made volatile.
220
221 class BogusFineUnionFind implements UnionFind {

```

```

222     private final Node[] nodes;
223
224     public BogusFineUnionFind(int count) {
225         this.nodes = new Node[count];
226         for (int x=0; x<count; x++)
227             nodes[x] = new Node(x);
228     }
229
230     public int find(int x) {
231         while (nodes[x].next != x)
232             x = nodes[x].next;
233         return x;
234     }
235
236     public void union(final int x, final int y) {
237         while (true) {
238             int rx = find(x), ry = find(y);
239             if (rx == ry)
240                 return;
241             synchronized (nodes[rx]) {
242                 synchronized (nodes[ry]) {
243                     // Check rx, ry are still roots, else restart
244                     if (nodes[rx].next != rx || nodes[ry].next != ry)
245                         continue;
246                     if (nodes[rx].rank > nodes[ry].rank) {
247                         int tmp = rx; rx = ry; ry = tmp;
248                     }
249                     // Now nodes[rx].rank <= nodes[ry].rank
250                     nodes[rx].next = ry;
251                     if (nodes[rx].rank == nodes[ry].rank)
252                         nodes[ry].rank++;
253                     compress(x, ry);
254                     compress(y, ry);
255                 }
256             }
257         }
258     }
259
260     // Assumes lock is held on nodes[root]
261     private void compress(int x, final int root) {
262         while (nodes[x].next != x) {
263             int next = nodes[x].next;
264             nodes[x].next = root;
265             x = next;
266         }
267     }
268
269     public boolean sameSet(int x, int y) {
270         return find(x) == find(y);
271     }
272
273     class Node {
274         private volatile int next, rank;

```

```
275
276     public Node(int next) {
277         this.next = next;
278     }
279 }
280 }
```

---

### 8.3 MyStack.java

---

```
1  import java.lang.*;
2  import java.util.*;
3  import java.util.concurrent.*;
4  import java.util.concurrent.atomic.*;
5  public class MyStack<T> {
6      private Object lock;
7      private final List<LinkedList<T>> stacks;
8      private static final int STRIPES = 32;
9
10     public MyStack(){
11         lock = new Object();
12         stacks = new ArrayList<LinkedList<T>>();
13         for(int i = 0; i < STRIPES; i++)
14             stacks.add(new LinkedList<T>());
15     }
16
17     public void push(T obj) {
18         int stripe = Thread.currentThread().hashCode()%STRIPES;
19         LinkedList<T> stack = stacks.get(stripe);
20         synchronized(stack){
21             stack.push(obj);
22         }
23     }
24
25     public T pop() {
26         int stripe = Thread.currentThread().hashCode()%STRIPES;
27         for (int i = stripe; i < stripe+STRIPES; i++){
28             LinkedList<T> stack = stacks.get(i%STRIPES);
29             synchronized(stack){
30                 if(stack.size() == 0) continue;
31                 return i == stripe ? stack.pop() : stack.removeLast();
32             }
33         }
34         return null;
35     }
36
37     public static void concurrentTest(final int size, final int threads,
38         MyStack<Integer> stack) throws Exception {
39         final CyclicBarrier startBarrier = new CyclicBarrier(threads+1),
40             stopBarrier = startBarrier;
41
42         final int range = size/threads;
```

```

43     for (int i = 0; i < threads; ++i) {
44         final int nr = i;
45         Thread ti = new Thread(new Runnable() { public void run() {
46             try { startBarrier.await(); } catch (Exception exn) { }
47                 for(int j = range*nr; j<range*nr+range; j++)
48                     stack.push(j);
49             try { stopBarrier.await(); } catch (Exception exn) { }
50             }});
51         ti.start();
52     }
53     startBarrier.await();
54     stopBarrier.await();
55     startBarrier.reset();
56     stopBarrier.reset();
57
58     final Set<Integer> pops = ConcurrentHashMap.newKeySet();
59     for (int i = 0; i < threads; ++i) {
60         final int nr = i;
61         Thread ti = new Thread(new Runnable() { public void run() {
62             try { startBarrier.await(); } catch (Exception exn) { }
63                 for(int j = range*nr; j<range*nr+range; j++)
64                     pops.add(stack.pop());
65             try { stopBarrier.await(); } catch (Exception exn) { }
66             }});
67         ti.start();
68     }
69
70     startBarrier.await();
71     stopBarrier.await();
72
73     if (pops.size() == size) System.out.println("Concurrency works :");
74     else System.out.println("Concurrency doesn't work :");
75 }
76
77 public static void testOrder(int n, MyStack<Integer> stack){
78     final AtomicBoolean working = new AtomicBoolean(true);
79     Thread A = new Thread(new Runnable() { public void run() {
80         for(int i=0; i<n; i++) stack.push(i);
81     }});
82     Thread B = new Thread(new Runnable() { public void run() {
83         for(int i=0; i<n; i++)
84             working.compareAndSet(true, n-1-i == stack.pop());
85     }});
86     try {
87         A.start(); A.join();
88         B.start(); B.join();
89     } catch (Exception e) {
90         System.out.println("Order dies >:");
91     }
92     if(working.get()) System.out.println("Order works :");
93     else System.out.println("Order doesn't work :");
94 }
95

```



```

96     public static void main(String[] args){
97         int size = 10_000_000;
98         int threads = 32;
99         MyStack<Integer> stack = new MyStack<Integer>();
100        testOrder(size, stack);
101        try {
102            concurrentTest(size, threads, stack);
103        } catch (Exception e){
104            System.out.println("Concurrent test died >:(");
105        }
106        System.exit(0);
107    }
108 }

```

---

## 8.4 SecComSys.java

---

```

1  // COMPILE:
2  // javac -cp scala.jar:akka-actor.jar SecComSys.java
3  // RUN:
4  // java -cp scala.jar:akka-actor.jar:akka-config.jar:. SecComSys
5
6  import java.util.*;
7  import java.io.*;
8  import akka.actor.*;
9
10 // -- HANDOUT -----
11 class KeyPair implements Serializable {
12     public final int public_key, private_key;
13     public KeyPair(int public_key, int private_key) {
14         this.public_key = public_key;
15         this.private_key = private_key;
16     }
17 }
18
19 class Crypto {
20     static KeyPair keygen() {
21         int public_key = (new Random()).nextInt(25)+1;
22         int private_key = 26 - public_key;
23         System.out.println("public key: " + public_key);
24         System.out.println("private key: " + private_key);
25         return new KeyPair(public_key, private_key);
26     }
27
28     static String encrypt(String cleartext, int key) {
29         StringBuffer encrypted = new StringBuffer();
30         for (int i=0; i<cleartext.length(); i++) {
31             encrypted.append((char) ('A' + (((int)
32                                     cleartext.charAt(i)) - 'A' + key) % 26)));
33         }
34         return "" + encrypted;
35     }

```

```

36 }
37
38 // -- MESSAGES -----
39
40 class InitMessage implements Serializable {
41     public final ActorRef R;
42     public InitMessage(ActorRef R) {
43         this.R = R;
44     }
45 }
46
47 class RegisterMessage implements Serializable {
48     public final ActorRef pid;
49     public RegisterMessage(ActorRef pid) {
50         this.pid = pid;
51     }
52 }
53
54 class LookupMessage implements Serializable {
55     public final ActorRef pid;
56     public final ActorRef returnTo;
57     public LookupMessage(ActorRef pid, ActorRef returnTo) {
58         this.pid = pid;
59         this.returnTo = returnTo;
60     }
61 }
62
63 class KeyPairMessage implements Serializable {
64     public final KeyPair keyPair;
65     public KeyPairMessage(KeyPair keyPair) {
66         this.keyPair = keyPair;
67     }
68 }
69
70 class Message implements Serializable {
71     public final String Y;
72     public Message(String Y) {
73         this.Y = Y;
74     }
75 }
76
77 class CommMessage implements Serializable {
78     public final ActorRef pid;
79     public CommMessage(ActorRef pid) {
80         this.pid = pid;
81     }
82 }
83
84 class PubKeyMessage implements Serializable {
85     public final ActorRef recipient;
86     public final Integer publicKey;
87     public PubKeyMessage(ActorRef recipient, int publicKey) {
88         this.recipient = recipient;

```

```

89         this.publicKey = publicKey;
90     }
91 }
92
93 // -- ACTORS -----
94
95 class RegistryActor extends UntypedActor {
96     public final Map<ActorRef, Integer> registry = new HashMap<>();
97
98     public void onReceive(Object o) throws Exception {
99         if (o instanceof RegisterMessage) {
100             RegisterMessage rm = (RegisterMessage) o;
101             KeyPair keyPair = Crypto.keygen();
102             registry.put(rm.pid, keyPair.public_key);
103             rm.pid.tell(new KeyPairMessage(keyPair), getSelf());
104         } else if (o instanceof LookupMessage) {
105             LookupMessage lm = (LookupMessage) o;
106             lm.returnTo.tell(new PubKeyMessage(lm.pid, registry.get(lm.pid)), getSelf());
107         }
108     }
109 }
110
111 class ReceiverActor extends UntypedActor {
112     public ActorRef registry;
113     public int publicKey;
114     public int privateKey;
115
116     public void onReceive(Object o) throws Exception {
117         if (o instanceof InitMessage) {
118             InitMessage im = (InitMessage) o;
119             im.R.tell(new RegisterMessage(getSelf()), getSelf());
120         } else if (o instanceof KeyPairMessage) {
121             KeyPairMessage kpm = (KeyPairMessage) o;
122             publicKey = kpm.keyPair.public_key;
123             privateKey = kpm.keyPair.private_key;
124         } else if (o instanceof Message) {
125             Message m = (Message) o;
126             String X = Crypto.encrypt(m.Y, privateKey);
127             System.out.print("decrypted: " + X + "\n");
128         }
129     }
130 }
131
132 class SenderActor extends UntypedActor {
133     public ActorRef registry;
134
135     public void onReceive(Object o) throws Exception {
136         if (o instanceof InitMessage) {
137             InitMessage im = (InitMessage) o;
138             registry = im.R;
139         } else if (o instanceof CommMessage) {
140             CommMessage cm = (CommMessage) o;
141             registry.tell(new LookupMessage(cm.pid, getSelf()), getSelf());

```

```

142     } else if (o instanceof PubKeyMessage) {
143         PubKeyMessage pkm = (PubKeyMessage) o;
144         String X = "SECRET";
145         System.out.print("cleartext: '" + X + "'\n");
146         String Y = Crypto.encrypt(X, pkm.publicKey);
147         System.out.print("encrypted: '" + Y + "'\n");
148         pkm.recipient.tell(new Message(Y), getSelf());
149     }
150 }
151 }
152 //
153 // -- MAIN -----
154
155 public class SecComSys {
156     public static void main(String[] args) {
157         final ActorSystem system = ActorSystem.create("SecComSys");
158         final ActorRef registry = system.actorOf(Props.create(RegistryActor.class),
159             "reigstry");
160         final ActorRef receiver = system.actorOf(Props.create(ReceiverActor.class),
161             "receiver");
162         receiver.tell(new InitMessage(registry), ActorRef.noSender());
163         final ActorRef sender = system.actorOf(Props.create(SenderActor.class),
164             "sender");
165         sender.tell(new InitMessage(registry), ActorRef.noSender());
166         sender.tell(new CommMessage(receiver), ActorRef.noSender());
167         system.shutdown();
168     }
169 }

```

---